

PRAVEEN GUPTA

Software Engineer — Backend, Full Stack & GenAI/LLM Systems

praweengupta11001@gmail.com | +91 63885 53560 | Varanasi, UP, India | github.com/PraveenGupta11001 | linkedin.com/in/praveen-gupta-b783791b4

PROFESSIONAL SUMMARY

Software engineer with 3 years of production experience across backend systems, full-stack SaaS delivery, and applied GenAI/LLM engineering. Designs and scales REST APIs and multi-tenant backends (Python, FastAPI, Django, DRF), ships React/Next.js front ends end-to-end, and builds LLM-powered features — RAG, agentic workflows (LangChain, LangGraph), and tool-calling — on top of that same production infrastructure. Delivered up to 80% API performance gains, built OAuth 2.0 integrations across a dozen+ third-party platforms, and independently productionized a full Azure stack (VMs, ACR, managed Postgres/Redis, Key Vault, Nginx, Cloudflare Tunnel, GitHub Actions CI/CD) alongside hands-on Azure AI Foundry work (GPT-4.1-mini, Foundry Agents, Azure AI Search). Comfortable owning a feature from architecture and data modelling through frontend UX, deployment, and monitoring — open to backend, full-stack, or GenAI-focused engineering roles.

CORE SKILLS

Backend: Python, FastAPI, Django, Django REST Framework, REST API Design, JWT, OAuth 2.0, Celery, Redis

Frontend: React.js, Next.js, TypeScript, JavaScript, Tailwind CSS, Shadcn UI

AI / GenAI: LangChain, LangGraph, RAG, AI Agents, Agentic Systems, LLM Workflows, Vector Databases, Prompt Engineering, Tool Calling, Azure AI Foundry (GPT-4.1-mini, Foundry Agents, Azure AI Search)

Databases: PostgreSQL, MySQL, Snowflake — schema design, indexing, query optimisation, vector-search-ready design

DevOps & Cloud: Docker, GitHub Actions (CI/CD), AWS, Azure (VMs, ACR, Key Vault, VNets), Azure Blob Storage, Azure Data Factory, Nginx, Cloudflare / Cloudflare Tunnel, Sentry

Integrations: Klaviyo, Wix, Apollo.io, Hugging Face, Salesforce, HubSpot, Pipedrive, Google Sheets, GitHub, Jira, Zapier, n8n, Webhook Automation

PROFESSIONAL EXPERIENCE

AI Engineer — Chirpn IT Solutions LLP *May 2026 – Jun 2026*

Remote

- Engineered agentic AI workflows and LLM-powered features on the Perceivnow.ai platform — prompt pipelines, tool-calling chains, and RAG modules that surface actionable intelligence from integrated data sources.
- Architected and maintained the OAuth 2.0 integration layer (FastAPI) feeding the AI system live data from six platforms — Klaviyo, Wix, Apollo.io, Hugging Face, GitHub, and Jira — with token-refresh and multi-tenant credential storage.
- Delivered the supporting full-stack surface (React, Next.js, Supabase Auth) and PostgreSQL schema so AI features shipped as complete product experiences.
- Owned solution design and technical scoping for new features end-to-end, from architecture review through implementation, QA, and production deployment.

Full Stack Developer — DraftssTech Technologies *Aug 2025 – Feb 2026*

Remote

- Engineered and scaled a multi-tenant SaaS platform end-to-end (Django, DRF, PostgreSQL, Redis, Celery, React/TypeScript), cutting average API response time by ~80% through query optimisation and caching.
- Built AI-assisted workflow automation, including Zapier connector plugins and n8n webhook orchestration pipelines, eliminating repetitive data-entry tasks for client teams.
- Designed and shipped OAuth 2.0 integrations with Salesforce, HubSpot, Google Sheets, Pipedrive, Zapier, and n8n, enabling automated two-way data sync for enterprise customers.
- Provisioned production AWS infrastructure (SNS, SES, CloudWatch, Route 53) with end-to-end Sentry observability, and collaborated cross-functionally via PR reviews and CI pipelines (GitHub Actions, Docker).

Software Developer — Hayara IT Services Pvt Ltd *Jun 2023 – Jul 2025*

Varanasi, UP

- Built RESTful microservices (FastAPI, PostgreSQL) with JWT/OAuth authentication, and developed full-stack SaaS features with React.js and Tailwind CSS from Figma design through production deployment.

- Reduced API latency by 40–60% through Redis caching and query optimisation; managed backend infrastructure on AWS S3 and Azure Blob Storage.
- Constructed end-to-end ETL pipelines (DBT, Snowflake, Azure Data Factory) — data-pipeline experience directly applicable to RAG ingestion and embedding workflows.

Software Engineer Intern — Logic Inventor Technologies *Oct 2022 – Mar 2023*

Varanasi, UP

- Developed and optimised REST APIs and backend business logic, building foundational API-engineering skills.

KEY PROJECTS

PayPilot AI — Enterprise Payroll, HR & Employee Self-Service Platform — payroll-saas-omega.vercel.app

- Designed a secure dual-role (Employer/Employee) authentication system with role-based access control and dashboard routing; integrated OAuth 2.0 PKCE with Slack and Jira.
- Built an async SMTP/OTP onboarding service (FastAPI BackgroundTasks) and achieved 100% unit test coverage on core auth modules (Pytest, Vitest, TDD).
- Used the platform as a hands-on Azure AI testbed — GPT-4.1-mini, Foundry Agents/tool-calling, and Azure AI Search vector retrieval — and productionized it on Azure (VMs, ACR, Key Vault, VNets, Nginx, Cloudflare Tunnel, GitHub Actions CI/CD); live at paypilot.llamateridon.site.

DocuChat — Production RAG Document Intelligence System — ai-doc-chat-flax.vercel.app

- Designed and shipped a document search and Q&A system (LangChain, FastAPI) with vector-based ingestion, chunking, and re-ranking for sub-second semantic search over large document corpora.

WeConnect — Real-Time Chat Platform with Pub/Sub Architecture — we-connect-teal.vercel.app

- Built a real-time messaging platform (FastAPI WebSockets, Redis Pub/Sub, React.js) with room-based routing and presence management — an event-driven pattern also used for streaming AI agent responses.

EDUCATION

Master of Computer Applications (MCA) — Arunachal University of Studies

Bachelor of Computer Applications (BCA) — Mahatma Gandhi Kashi Vidyapeeth University

LANGUAGES

Hindi (Native) · English (Professional)